

pgrep: A Physics GRE Preparation System with Separated, Calibrated Readiness Scores

Frank Gonzalez

July 2026

Abstract

pgrep is a desktop and mobile study system for the Physics GRE (PGRE). The PGRE rewards problem solving over fact recall, so a recall-only flashcard tool is the wrong instrument for it. pgrep addresses this with four pieces of work. First, an in-engine study selector reorders due items by exam value, blueprint weight times topic weakness, without altering any scheduling state. Second, three separated and calibrated scores report what a learner can recall now (Memory), whether they can apply it to a new question (Performance), and what they would likely score (Readiness), each with an explicit range and an abstain rule. Third, an AI item layer generates cards, misconception-first multiple-choice problems, and a scaffold-fade tutor, each citing a named open source or refusing, and each gated by a pre-registered gold-set test and a leakage firewall. Fourth, every model is evaluated on held-out data against a bar set in advance, and the negatives are reported. On held-out review logs, Memory calibrates to a Brier score of 0.234 and beats a base-rate baseline. A simulated ablation shows the selector beats massed practice in every configuration and beats stock scheduling at realistic study budgets. The AI layer beats keyword, vector, and naive-distractor baselines with confidence intervals excluding zero, although two absolute quality cutoffs remain just short. pgrep runs and scores with AI switched off.

Keywords: spaced repetition, exam readiness, calibration, interleaving, educational data mining, retrieval-augmented generation, item generation.

1 Introduction

Preparing for a single high-stakes exam is a measurement problem as much as a study problem. A learner needs to know not only whether a fact is in memory, but whether they can deploy it on a novel question under time pressure, and what score that predicts. The Physics GRE is a sharp instance. It is about seventy five-choice multiple-choice questions in two hours with no penalty for guessing, and its content blueprint has been stable for two decades: Mechanics 20%, Electromagnetism 18%, Quantum 13%, and six smaller areas. The exam is problem-solving-heavy and fact-recall-light. That single fact reframes the tooling. The intended user is a post-undergraduate physics student who has already met the material and wants to raise a score, not a novice meeting it for the first time, which tilts the design toward practice and measurement and away from teaching. Flashcards, the default study technology, train recall, and recall is the weakest predictor of the skill the PGRE actually tests.

The gap between the two is not incidental. Soderstrom and Bjork (2015) distinguish *learning*, a durable and transferable change, from *performance*, what a learner can do right now, and show that the two often move in opposite directions. A study method can inflate in-session performance while

doing little for durable transfer. A study tool that measures only recall therefore reports a number that feels like readiness but is not. pgreg is built on the premise that three different questions need three different instruments. Can you recall a fact now (Memory)? Can you apply it to a question you have not seen (Performance)? What would you score (Readiness)? These are kept as three separate numbers, never collapsed, because recall does not imply transfer. A learner can have strong Memory on a topic and weak Performance on it, having memorized a result without being able to choose it on a fresh problem, and folding the two into one number would hide exactly the gap the exam probes.

The second premise is that honesty is a design property, not a disclaimer. Every number pgreg shows carries a likely range, a coverage figure, and an abstain rule that refuses to render a score when the data behind it is too thin. A readiness estimate over a topic the learner has never practiced is not a low number, it is no number, and the surface says which topics are missing. This posture, that abstaining beats bluffing, runs through the whole system and through this report.

pgrep is a working desktop application with a native mobile companion, sharing one study engine with two-way sync. It reuses a proven memory model for scheduling and builds the exam-readiness layer on top. This report describes the system and evaluates it. The contributions are four.

1. **An in-engine exam-value selector.** A new review-card ordering, computed inside the study engine, reorders due items by blueprint weight times topic weakness, within a desirable-difficulty band, without mutating any scheduling state.
2. **Three separated, calibrated scores.** Memory, Performance, and Readiness, each with an 80% central interval and an abstain rule, and each validated against held-out outcomes rather than asserted.
3. **A provenance-gated AI item layer.** Card generation, misconception-first problem generation, and a scaffold-fade tutor, each of which cites a named open-corpus source or refuses, and each gated by a pre-registered gold-set test and a leakage firewall.
4. **A reproducible held-out evaluation.** Every model is measured on held-out data against a bar fixed before results were seen, and the negatives are reported alongside the wins.

The remainder of the report covers the foundation pgreg reuses, the system at a glance, the study selector, the three-score model, the AI layer, the evaluation, and a discussion of what the evidence does and does not support.

2 Background and Foundation

2.1 The reused engine

pgrep is built by reusing the engine of an open-source spaced-repetition application (Anki) and replacing its interface and its study logic. The reuse is deliberate and bounded. What pgreg inherits is the memory model, the collection store and its review log, the two-way sync system, and the local web-host mechanism that serves the user interface. What pgreg adds, and what this report is about, is everything above that line: the exam-value selector, the three-score model, the AI item layer, the evaluation harness, the desktop surfaces, and the mobile companion. The reused engine is licensed under the AGPL, and pgreg preserves that license and its attribution.

The memory model is worth naming because the Memory score depends on it. The engine ships FSRS, a modern three-component model of memory (difficulty, stability, retrievability) whose forgetting curve gives, for any card at any moment, a predicted probability of recall. On a large public benchmark of review histories, FSRS is the strongest non-neural scheduler, ahead of the older SM-2 heuristic and of half-life regression (Settles & Meeder, 2016), and it descends from a line of work that frames spaced repetition as an optimization problem (Open Spaced Repetition, 2024b; Tabibian et al., 2019; Ye, Su, & Cao, 2022). pgreg does not modify the memory model. It consumes the model’s retrievability estimate and builds selection and scoring on top of it.

2.2 The learning-science basis

Four findings from the learning-science literature shape pgreg’s features. Each is a mechanism pgreg implements, not a result pgreg claims for itself.

Interleaving. Practicing related topics in a mixed order, rather than in blocks, improves later performance on novel problems, often substantially. With spacing held fixed, interleaved mathematics practice roughly doubled next-day test scores (Taylor & Rohrer, 2010), and the benefit extends to problems that are not superficially similar, with large effects reported, a Cohen’s d near 1.0 in one multi-week study (Rohrer, Dedrick, & Burgess, 2014; Rohrer & Taylor, 2007). In undergraduate physics, mixed homework raised surprise novel-problem scores by 50 to 125 percent (Samani & Pan, 2021), and meta-analyses report medium effects that are largest when the concepts are confusable (Firth, Rivers, & Boyle, 2021; Sana & Yan, 2022). The mechanism matters here: interleaving forces the learner to choose a strategy from the problem itself, training the discrimination the PGRE demands. Learners reliably misjudge its value, preferring blocked study that feels easier (Kornell & Bjork, 2008).

The generation effect. Material a learner produces is remembered better than material they read, even when the content is identical, with a meta-analytic effect around $d = 0.40$ (Bertsch, Pesta, Wiscott, & McDaniel, 2007; Slamecka & Graf, 1978). User-generated flashcards outperform premade ones on both memory and application (Pan, Zung, Imundo, Zhang, & Qiu, 2022). pgreg uses this to gate AI card generation behind an authoring step, so the learner pays the generative cost at novel-concept formation before the system amortizes that style across the rest of the deck.

Productive struggle with consolidation. Attempting a problem before receiving help, then consolidating with feedback, aids transfer (Kapur, 2008, 2016), subject to the caveat that heavy guidance backfires as prior knowledge grows (Kalyuga, 2007). In a randomized trial with about one thousand students, unguarded access to a general assistant reduced learning, whereas a guardrailed tutor did not (Bastani et al., 2024). Help must therefore be scaffolded, must make the learner produce the reasoning (Chi, Bassok, Lewis, Reimann, & Glaser, 1989), and must not leak the answer, since a large share of naive AI hints simply solve the task for the student. Retrieving principles before solving primes the solving itself, with an effect near $d = 0.4$ (Gjerde et al., 2022; Roediger & Karpicke, 2006), and completion and worked-example structure manage cognitive load (Renkl & Atkinson, 2003; Sweller, 1988).

Calibration over confidence. The honest move is to make the system’s own predictions calibrated and visible, not to harvest the learner’s subjective confidence. This is measured with proper scores and reliability diagrams (Brier, 1950; DeGroot & Fienberg, 1983; Guo, Pleiss, Sun, & Weinberger, 2017).




Figure 1: System data flow. The attempt log and review log are the spine: every card review, problem attempt, and exam item feeds them, they drive the three scores and the honest dashboard, and the FSRS state feeds both the Memory score and the points-at-stake selector that orders the study doors.

3 System Overview

pgrep is organized around three co-equal pillars that map onto the three scores, and a single study loop that feeds them (Figure 1).

The three pillars. Cards build Memory through retrieval. Problems build Performance through application, with reasoning-first help. Timed mocks build Readiness under exam conditions. The pillars are kept co-equal on purpose. The thesis that the PGRE rewards problem solving is not enforced by a hand-tuned weighting, it falls out of the engine and the study mix, so no pillar is special-cased.

The study loop. A session runs as two doors, never one shuffled queue, because mixing cards and problems within a single queue has no support in the literature. Inside each door, topics are interleaved. On the Problems door, the learner commits an answer before any help is offered, which preserves the attempt-before-help structure that productive struggle requires. A short adaptive diagnostic places each topic as strong or rusty at first run, and a coverage-aware daily mix suggests the next best thing to study. A Focus drill reuses both doors scoped to a single topic for targeted repair, and every Problems session ends with a short synthesis of the patterns across its problems, the consolidation step that productive struggle requires. Exam mode is kept separate as the measuring instrument, timed and help-free, so the number used to estimate readiness is never inflated by the help that makes everyday practice productive.

Surfaces and platforms. The desktop application presents six custom surfaces: Home (the three scores at a glance), Study (the two doors plus a timed exam mode), Progress (coverage and model calibration), Diagnostic, Library (authoring and generation), and Settings. A native mobile

companion mirrors Home and Study on the same shared engine, and the two devices sync two-way. The desktop and mobile hosts are thin shells over one engine and one remote-procedure surface, so parity is structural rather than duplicated.

The attempt log is the spine. Every graded event, a card review, a problem attempt, an exam item, is recorded. Card reviews ride the engine’s existing review log. Problem attempts, which carry richer data (the selected option, the depth reached in the help ladder, sub-goal productions, timing, and a session id), are stored as an append-only log of immutable notes. This choice, storing the log as notes rather than as a custom database table, lets it ride the engine’s existing sync for free and merge cleanly across devices, at the cost of a little analytic overhead. The log is the single source that feeds the selector, the three scores, and the calibration dashboard.

3.1 Architecture and data model

Desktop and mobile are two thin hosts over one engine. On the desktop, a local web server serves the interface, a native window displays it, and the page calls the engine through the same remote procedures the mobile app uses. The mobile companion is a native application that drives the shared engine through a small foreign-function-interface layer. Because both hosts call one engine and one procedure surface, including the exam-value selector, parity between devices is structural rather than duplicated, and a study session behaves the same on either.

The attempt log’s storage was a deliberate decision with a documented evolution path, recorded as *A now, C-ready*. Storing each attempt as an immutable note (option A) lets the log ride the engine’s note sync for free and merge across devices by identifier, at the cost of parsing string fields for analytics. A custom database table (option B) would give faster typed queries but would require writing custom sync inside the engine, the project’s most heavily weighted requirement, so it was rejected. A local recomputable cache folded from the notes (option C) recovers that speed with no sync risk, and is pre-engineered but deferred until the fold measurably slows. Five invariants make the later flip to C free: notes are immutable and self-contained, event identity is the note’s own identifier, the payload is one JSON blob plus a few hot fields, all analytics pass through a single read seam, and the cache is a pure fold that never touches sync.

4 The Study Selector

The first pillar’s engine change is a new way to order the cards that are due on a given day. The design goal is simple to state: when there are more due cards than the daily limit allows, spend the limit on the cards that matter most for the exam, and present them in an interleaved order.

4.1 Why this is a real engine change

A naive implementation would gather the day’s due cards in some default order and re-sort them in the interface. That does not work here, because the engine applies the daily review limit *during* the database query that gathers cards, ordered by a fixed clause. The cards that matter most might never be gathered before the limit cuts the list off. Ordering by exam value therefore requires a change inside the engine, not a cosmetic sort after the fact.

pgrep adds a new review-card ordering, `ReviewCardOrder::PointsAtStake`, and a second pass in the engine that gathers all due reviews, scores them, and then truncates to the limit. This




Figure 2: The points-at-stake selector as a gather-then-limit second pass. It gathers every due review, scores each by exam value within a difficulty band and with anti-blocking, and truncates to the daily limit, all in memory, so no scheduling state is written.

gather-then-limit structure is the substantive change (Figure 2).

4.2 The score

Each due card is scored by the exam value at stake in reviewing it:

$$\text{worth}(\text{card}) = \text{blueprint\%}(\text{topic}) \times \text{weakness}(\text{topic}), \quad \text{weakness}(\text{topic}) = 1 - \bar{R}_{\text{topic}},$$

where $\bar{R}_{\text{topic}}$ is the mean predicted retrievability over that topic’s due cards, and `blueprint%` is the official PGRE weight for the topic’s area. Weakness is computed from the memory model alone, so the selector needs no attempt-log dependency and the per-topic weakness map can be cached once per session. Among the scored cards, the selector prefers those in a desirable-difficulty band of roughly 60 to 85% predicted success (Metcalfe & Kornell, 2005; Wilson, Shenhav, Straccia, & Cohen, 2019), and it applies an anti-blocking rule that allows at most three consecutive cards from the same area, so the emitted order stays interleaved.

Computing weakness from the memory model has a practical payoff. The per-topic weakness map is built once per session and cached, so the second pass stays well under the interface’s responsiveness budget even on a large collection. The shipped weighting is fixed, which is the graded minimum. A planned extension anneals the balance between worth and the difficulty band along a dial driven by time to test and coverage, shifting from a learning emphasis when the exam is far toward a triage emphasis as it nears. That loop steers study only, never the held-out measurement, so it can never grade its own homework.

4.3 The safe seam

The selector reorders the in-memory set of due cards and nothing more. It never mutates the scheduling fields (`due`, `interval`, or `memory_state`), never writes collection data, and never creates an undo entry. Queues live in memory and are rebuilt each session, so reordering them writes nothing and cannot corrupt the collection. This is what lets a genuine engine change sit safely inside a scheduler the authors did not write: scheduling, undo, and sync are untouched, and the memory model’s validity is preserved by construction. The change is covered by three engine-level tests and one integration test.

Score	Predicts	Model	Abstains when
Memory	Recall now	Blueprint-weighted FSRS retrievability	A topic has fewer than 5 re- viewed cards
Performance	Correct on a new ques- tion	PFA calibrated logistic over 4 features	A topic has fewer than about 8 clean attempts
Readiness	Projected scaled score	Expected raw mapped to the 200–990 scale	Coverage is under 70% of blueprint weight

Table 1: The three scores. Each is a distinct construct, model, and abstain rule.




Figure 3: How the three scores are derived. Memory is a blueprint-weighted average of FSRS retrievability; Performance is a calibrated logistic over attempt-log features; Readiness maps the expected raw score, built from per-topic p_t and blueprint question counts n_t , through the official conversion table. Each carries an 80% range and an abstain rule.

5 The Three-Score Model

The core of pgrep is three separate numbers, each answering a different question, each reported with a range and an abstain rule. The design principle is that uncertainty is first-class and standardized: every range is an 80% central interval, computed the same way everywhere, and every score refuses to render when its data is too thin. Three further principles hold across all three. Each number carries its own honesty, a point estimate, a likely range, a data-sufficiency read, and a last-updated time, so no bare number appears without its evidence. All three are pure arithmetic over the memory model and the attempt log, with no model in the loop, so both apps produce all three scores with AI switched off. And each reuses the engine’s own primitives where it can, so Memory and the selector’s weakness term, for one, are the same quantity and cannot disagree. Table 1 summarizes the three.

5.1 Memory

Memory is the expected fraction of a learner’s due and known cards they would recall right now, weighted by the blueprint. It reuses the memory model directly: each card’s predicted retrievability is read from the engine, averaged per topic, and combined across topics by blueprint weight. Because the number of cards recalled is a sum of independent, non-identical Bernoulli trials, its distribution is Poisson-binomial (Le Cam, 1960), which gives an analytic mean and variance and hence the 80% range. A topic with fewer than five reviewed cards shows “not enough cards yet” rather than a number. By construction, Memory and the selector’s weakness term use the same primitive, so the score and the study order never disagree.

5.2 Performance

Performance is the probability of answering a new, unseen exam-style question correctly, per topic and overall. This is transfer, not recall, so it is computed from the attempt log rather than from the memory model. The model is Performance Factors Analysis, a calibrated logistic regression that is the standard education-data-mining approach to predicting next-item correctness (Cen, Koedinger, & Junker, 2006; Pavlik, Cen, & Koedinger, 2009). It is preferred over Bayesian Knowledge Tracing, which has known identifiability difficulties (Corbett & Anderson, 1994), and over deep knowledge tracing, which is data-hungry and opaque at this scale (Piech et al., 2015). For a problem j in topic t ,

$$P(\text{correct}) = \sigma(\beta_0 + \beta_m m_t - \beta_d d_j + \gamma_s s_t + \gamma_f f_t),$$

where m_t is topic mastery (mean retrievability), d_j is the authored item difficulty, and s_t and f_t are recency-weighted counts of recent successes and failures, kept separate because successes and failures carry different weight (Galyardt & Goldin, 2014). Response latency is used only to filter rapid-guess attempts and to measure pace in the timed exam mode, not as a predictor, because in untimed practice a slow answer is ambiguous. Coverage is used to widen the interval and to gate Readiness, not as a predictor.

Two modeling choices are worth defending. First, the baseline the model must beat is the per-topic base rate, the batting average, which is honest and robust and is what the system falls back to with too little data. PFA earns its place only if it beats that baseline (Gong, Beck, & Heffernan, 2011). Second, a full item-response or Rasch model was rejected. With a single learner, item difficulty and learner ability are jointly unidentifiable, and such models need roughly one to two hundred examinees per item (Embretson & Reise, 2000; Lord, 1980). The authored difficulty tag gives most of the difficulty-awareness without the estimation. Raw logistic scores are not automatically calibrated, so a post-hoc beta calibration is fit on a held-out split (Kull, Silva Filho, & Flach, 2017; Platt, 1999), chosen over alternatives that overfit at small sample sizes. The 80% interval comes from Bayesian partial pooling across topics (Gelman et al., 2013), and a topic abstains until its interval is tight enough, with about eight to ten clean attempts as a starting proxy (Brown, Cai, & DasGupta, 2001; Peduzzi, Concato, Kemper, Holford, & Feinstein, 1996).

5.3 Readiness

Readiness is the projected scaled score, in the PGRE’s 200 to 990 band, with an explicit range and a coverage figure. It leans on Performance, not Memory, because the exam measures transfer. For each area, the per-topic probability p_t from the Performance model is combined with n_t , the

number of exam questions that area contributes under the blueprint. The expected raw score is $\sum_t n_t p_t$, again a Poisson-binomial sum with an analytic mean and variance. The raw point and the raw interval endpoints are then mapped through the official raw-to-scaled conversion table from a real practice form, used as numeric constants only. The result is a scaled point estimate and an 80% scaled interval. Thin coverage widens the interval before it trips the gate, so the range tightens as more of the exam is practiced, which is the honest visual of readiness accruing.

Readiness is where honesty-by-construction is most visible. Coverage is the fraction of blueprint weight whose topics have cleared the Performance data floor. If coverage is under 70%, Readiness abstains, states that not enough of the exam is covered yet, and names the uncovered areas. A learner cannot manufacture a readiness number over a part of the exam they have not practiced.

6 The AI Layer

The AI layer enriches the item pool. It is an upgrade, never a dependency: every path has a working AI-off fallback, and an AI-off session loads none of the AI code. The layer has one hard rule, that every generated output cites a named source from an open corpus or is refused, and it is gated before anything ships.

6.1 What it generates

Cards. Card generation is pay-to-play, reflecting the generation effect. The learner authors a conceptual seed first. If the bundled deck already covers that subtopic, the AI stylizes the verified card into the learner’s voice, with the facts locked one-to-one. If the deck lacks the technique, the seed triggers gap-fill generation of new cards, grounded in the corpus and gold-set gated. Import adds breadth but never substitutes for authoring, because importing someone else’s cards yields no generation effect. Figure 4 shows the routing.

Problems. The value and the risk of a multiple-choice problem sit in its distractors. Naive generation asks the model for wrong answers and gets implausible ones. `pgrep` uses misconception-first generation: name the specific error a student is likely to make (a sign error, a wrong law, a unit slip, a limiting-case confusion), then derive the distractor that error actually produces, and store the misconception tag and rationale with each choice. Every generated problem also carries a solution decomposition, ordered sub-goals with a rubric each, verified once at creation. A problem without a verified decomposition does not ship. The misconception tags come from a small taxonomy of recurring physics errors, and the stored rationale for each distractor is what the tutor later teaches from, so the same annotation that makes a problem hard also makes its help specific (Figure 5).

The tutor. When a learner misses a problem, a scaffold-fade ladder helps without leaking the answer: a nudge, then sub-goal decomposition with self-explanation over the stored decomposition, then a sibling worked example, then a full reveal with explain-back. The ladder walks the stored, pre-verified sub-goals rather than generating steps live, which makes it fast, reliable, and AI-off capable. A giveaway verifier checks every AI hint and refuses rather than reveal the answer, addressing the leakage that unguarded assistance introduces (Bastani et al., 2024). With AI off, the same ladder runs as reveal-and-self-compare. Figure 6 shows the ladder.




Figure 4: Card generation is pay-to-play. The learner authors a seed first; with AI on, the system stylizes existing bundle cards where the deck already covers the subtopic and gap-fills new, corpus-grounded cards where it does not. Everything passes verification and the gold-set gate before entering the one shared pool.

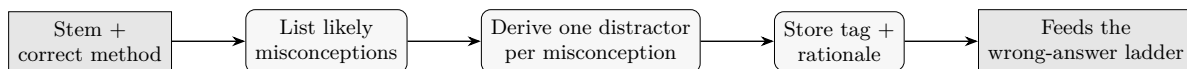


Figure 5: Misconception-first problem generation. Rather than asking for wrong answers, the pipeline names the likely error, derives the distractor that error produces, and stores the tag and rationale, which the tutor then teaches from.

6.2 Provenance, the gold-set gate, and the firewall

Three mechanisms keep the layer honest.

Provenance. Generation is grounded by retrieval over an open corpus (OpenStax University Physics and the Fitzpatrick texts), chunked and embedded, each chunk carrying a stable source reference. A generated item that cannot ground its facts in the corpus is refused.

The gold-set gate. Generation is graded against hand-verified gold sets with cutoffs fixed before any results are seen. For cards, fact precision must reach 0.95 and useful-yield 0.80. For problems, key correctness must reach 0.95, per-problem distractor quality 0.70, and useful-yield 0.75. The AI must also beat the better of a keyword and a vector retrieval baseline by at least 0.10 absolute on the headline metric, with the advantage’s confidence interval excluding zero. Two blind raters score every batch, a human and an AI judge, with the human adjudicating.

The leakage firewall. Held-out and gold items never enter the corpus, the retrieval index, or any prompt. The index is built from the corpus folder only, which is structural rather than merely




Figure 6: The wrong-answer ladder. Help fades in from a nudge to a full reveal, the learner produces reasoning at each rung over a stored, verified decomposition, and a giveaway verifier keeps the final answer from leaking. Any correct re-attempt exits to consolidation.

promised, and an automated guard asserts that no gold or held-out path and no held-out verbatim span appears in the index. A dedup scan found zero verbatim reprints of real exam items among the fed examples.

6.3 The verification stack and the data

Generation is verified in layers, ordered by how much each can be trusted. Retrieval grounding with a source anchor is the first and load-bearing layer, since it is what makes provenance mechanical. For computational items, a symbolic check with a computer-algebra system is decisive, and the pipeline routes items by kind: computational cards and problems are checkable and high-trust, while conceptual items lean on provenance and a human spot-check, which is where the residual risk concentrates. Self-consistency across multiple samples adds signal, and an independent-critic pass is treated as weak, because self-critics tend to rubber-stamp. Items below a grounding-confidence floor are routed to human review rather than shipped. Figure 7 traces one item through the pipeline.

The data behind the layer is split cleanly to keep the evaluation honest. The corpus the AI reads and cites is open, OpenStax University Physics and the Fitzpatrick texts, chunked and embedded into




Figure 7: The generation and verification pipeline for one item. Retrieval grounding gives provenance, a symbolic check verifies computational items while conceptual items lean on provenance and a spot-check, low-confidence items go to human review, and a gold-set gate governs what ships.

about four thousand four hundred passages, each with a stable source reference, and the keyword and vector baselines run off the same index. Real exam forms are never fed to generation. One older form is vision-cleaned into the problem gold ruler, two clean forms are held out as the performance bank and played in the app’s exam mode, two more are reserve held-out, and one is a sealed final mock that also supplies the raw-to-scaled constants. Few-shot examples come from a separate clean non-exam pool. Because nothing is both fed and graded, the beats-a-baseline-on-held-out claim is airtight, and the dedup scan confirmed no fed example restates a real exam item. Figure 8 shows the split.

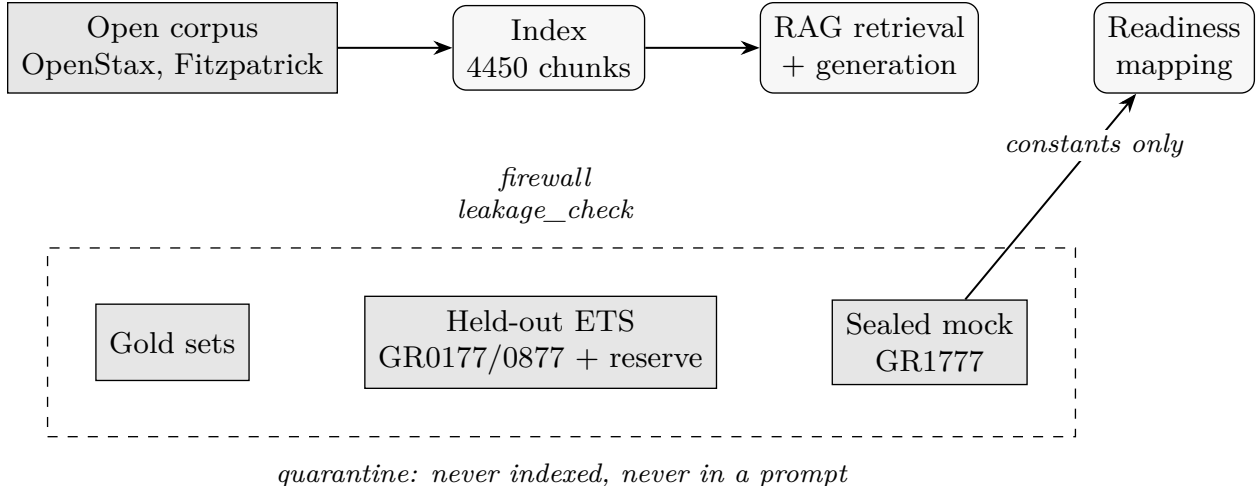


Figure 8: The leakage firewall. Only the open corpus is indexed and reachable by generation. The gold sets, the held-out exam forms, and the sealed mock stay quarantined. The one permitted crossing is the sealed mock’s numeric score-conversion constants, used by the Readiness mapping.

7 Evaluation

Every model is evaluated on held-out data against a bar set in advance, and the results are reported with their negatives. This section states the methodology, then the results for each component. The honest boundary throughout is that this is a single-learner project, so where a human cohort would be needed, the evaluation validates the mechanism or the pipeline and says so. Figure 9 shows the three held-out pipelines, one per score, that produce the numbers in this section.

7.1 Methodology

Three principles govern the evaluation. Splits are time-based, never random, so no information leaks across a card’s or an item’s trajectory (Kapoor & Narayanan, 2023). Cutoffs and baselines are pre-registered, frozen and dated before results are seen. Runs are reproducible from one command with fixed seeds, and every metric is reported with a bootstrap confidence interval. The primary calibration metric is the Brier score, which is binning-free, supported by log-loss, the expected calibration error, and a reliability diagram, which can disagree at small samples and so are reported together.

The scoring process is built to keep the numbers honest. Items from every system, the AI and each baseline, are shuffled together and scored blind to their source, so a rater cannot favor the AI. Two raters score each batch, a human and an AI judge, the human adjudicates disagreements, and inter-rater agreement is reported alongside the scores. Every run is driven by one command with fixed seeds and pinned splits, and each headline number carries a bootstrap confidence interval, so the whole evaluation reproduces from the same inputs.

7.2 Memory calibration

Memory is calibrated on a held-out tail of a large public review-log dataset (Open Spaced Repetition, 2024a), where the model predicts retrievability and the log records the actual pass or fail. On the




Figure 9: The held-out evaluation, one pipeline per score. All splits are time-based, all cutoffs and baselines are pre-registered, and one command reproduces every number with bootstrap confidence intervals.

held-out split, Memory reaches a Brier score of 0.234, a log-loss of 0.743, and an expected calibration error of 0.159, and it beats the base-rate baseline on the primary Brier score. The reconstruction is pinned to the exact memory-model version the engine ships, so the calibration reflects the deployed model.

7.3 Performance

With one learner there is not yet a bank of the learner’s own attempts on novel items large enough to fit and validate a logistic model honestly. The Performance pipeline is therefore validated on a seeded, documented synthetic dataset whose outcomes are drawn from a known logistic process, which validates the fitting, calibration, and held-out machinery rather than a human effect. On the held-out synthetic split, the calibrated model reaches a Brier score of 0.175 and an accuracy of 0.775, against 0.268 and 0.563 for the base-rate baseline, and the pre-registered beat-the-baseline rule passes. The report states plainly that this validates the pipeline at a sample size of one, and that real coefficients need a real cohort.

7.4 Readiness

Readiness maps expected performance to the 200 to 990 scale with an 80% range and is coverage-gated at 70%. With one or two practice forms available, the check of predicted against actual scaled score is a sanity check, reported as such rather than as a full validation. The mechanism of interest, the coverage gate and the abstain behavior, is the part that is verifiable now, and it behaves as specified: below the coverage line, Readiness abstains and names the uncovered exam.

7.5 The AI layer

The AI gate was scored on a batch against 157 verified gold items, with generation from a pinned model and two blind raters. Table 2 includes the headline numbers. Under the human adjudicator, generated cards reach a useful-yield of 0.84, clearing the 0.80 bar, with a fact precision of 0.90, one step under the 0.95 floor because five cards carried a fact slip. The non-refused problems are excellent by the human rater (key correctness 1.00, useful-yield 0.92, distractor quality 0.96), but a 31% problem refusal rate, mostly malformed or declined generations plus a smaller share of answer-leak safeguards, drags the batch numbers under the bars. The AI beats keyword and vector retrieval by 0.74 on cards and 0.67 on problems, and beats naive-distractor generation by 0.42, all with confidence intervals excluding zero, which is the spec’s core requirement and validates the misconception-first design.

Two findings deserve emphasis. First, the two absolute quality cutoffs that remain short (card fact precision and the problem batch yield) are generation issues, cutting the malformed-problem rate and tightening fact precision, not gold-set or methodology problems. Second, the human and AI judges barely agreed, with a Cohen’s kappa of about -0.03 on usefulness. The AI judge systematically under-rated usefulness, so an AI-only gate would have been unreliable here. That near-zero agreement is itself a result, and it is the clearest argument for the two-rater rule.

Two process details support these numbers. Between an initial pass and the scored run, the generator gained key self-consistency, an independent re-solve that regenerates a problem when its own solution disagrees with the key, and a tuned grounding floor, which together roughly halved refusals and lifted problem key-correctness under the AI judge. The firewall held throughout: no exam form was fed, and a reject-memorized check kept all eighty-six candidate items, none a near-copy of a fed example, with the seen-versus-held split recorded in the run manifest.

7.6 The interleaving ablation

The study feature is ablation-tested in a controlled simulation. With one learner there is no cohort, so the ablation validates the shipped selection mechanism under the memory model as ground truth, and does not establish a human interleaving effect. The pre-stated primary metric is blueprint-weighted expected recall at the exam horizon, which is the app’s actual objective and is coverage-aware. Three arms study for equal time: the full interleaved selector, a blocked (massed) arm, and stock scheduling in soonest-due order. Table 3 gives the grid across study budget and exam distance.

Two results stand out, and one non-result. The full selector beats massed practice in every configuration, by 0.008 to 0.075, with every confidence interval excluding zero. This is the core learning-science claim, that spaced topic-mixed practice beats massed practice, and it holds robustly. Against stock scheduling, the full selector wins at every realistic study budget, 20 and 30 reviews

Model	Metric (held-out)	Result	Baseline and note
Memory	Brier / log-loss / ECE	0.234 / 0.743 / 0.159	Beats base-rate on Brier; real review logs
Performance	Brier / accuracy	0.175 / 0.775	0.268 / 0.563 base-rate; synthetic, n=1 pipeline check
Readiness	Scaled 200–990, 80% range	Coverage-gated	Sanity check only, low n
AI cards	Useful-yield / fact precision	0.84 / 0.90	+0.74 vs retrieval; fact precision one step under bar
AI problems	Key / useful / distractor	1.00 / 0.92 / 0.96	Non-refused; +0.67 vs retrieval, +0.42 vs naive; 31% batch refusal

Table 2: Evaluation results. All comparisons are held-out; confidence intervals for the AI advantages exclude zero.

Budget (reviews/day)	Exam gap (days)	full – blocked	full – stock
10	0	+0.011	−0.037
10	90	+0.008	−0.023
20	0	+0.068	+0.005
20	90	+0.062	+0.005
30	0	+0.075	+0.015
30	90	+0.074	+0.014

Table 3: Ablation grid. Positive means the full selector is better. Every confidence interval excludes zero. The full selector beats massed practice everywhere and beats stock scheduling at realistic budgets, losing only at the tightest budget.

per day, including the headline configuration where it leads by 0.015. It loses only at the tightest budget, 10 reviews per day, and the reason is mechanical and expected: the desirable-difficulty band reviews cards below their soonest-due point, spending near-term recall to buy long-term stability, a trade that does not pay back when the budget is too small to review much at all. The non-result is that the anti-blocking reshuffle registers as exactly zero under this memory model, because same-day review order does not change the model’s updates. Its justification is cognitive, about discrimination and contextual interference, which this simulation cannot capture, so it is neither credited nor refuted here.

The comparison is controlled rather than casual. Every arm studies for exactly the same number of reviews, common random numbers pair the recall outcomes across arms so the arms differ only in their ordering, and the intervals are bootstrapped over forty seeds. The honest read of the blocked contrast is that it tests spaced-mixed versus massed practice, because the full arm also prioritizes within a topic by weakness, so the win over blocked is the broad interleaving effect rather than a single-variable swap. At the headline budget the full arm’s per-area recall is uniformly high across all nine blueprint areas, which is what a coverage-aware objective should produce.

7.7 Two-way sync

The two-device requirement is met by reusing the engine’s own sync unmodified against a self-hosted server. Reviews and problem attempts union by identifier, so two devices reviewing different cards

offline both land with nothing lost or double-counted, and for the same card on both devices the newer modification time wins the scheduling state. A round-trip test covers different cards offline, the same card on both, the attempt-log union, and offline-then-sync, and an on-device path uploads a review from the phone that the desktop then downloads. The conflict rule is the engine’s own, made explicit rather than reinvented: review history and the attempt log are append-only and union by identifier, so nothing is lost or double-counted, while for the same card edited on both devices the newer modification time wins the scheduling state, with a deterministic device tie-break on equal times. Offline safety rides the engine’s transactional store, so a session works fully offline and syncs opportunistically, and a kill-mid-review stress test leaves the collection uncorrupted.

8 Discussion

The through-line of pgreg is honesty by construction. The coverage gate and the abstain rules mean the system cannot present a readiness number over an unpracticed part of the exam, so a confident-looking but empty score is not a risk to be avoided by discipline, it is impossible by design. Every score ships a range and states its evidence, which is the practical form of preferring calibration to confidence.

The ablation illustrates the same posture applied to the authors’ own feature. The selector’s clean, robust win is over massed practice, the mechanism the literature actually supports. Against stock scheduling the win is real at realistic budgets but not universal, and the one loss is explained rather than hidden. A study tool that spends near-term recall to buy durable stability should be expected to lose when the learner can barely review, and it does. Reporting that makes the wins credible.

The AI evaluation produced a methodological result the project did not set out to find. The human and AI judges agreed at close to chance, and the AI judge under-rated usefulness. Had the gate relied on the AI judge alone, it would have rejected good material. The two-rater rule, adopted in advance for integrity, turned out to be load-bearing. This is a small but concrete case for keeping a human in the loop on generative-quality gates.

Finally, AI is an upgrade and not a dependency. The three scores are pure arithmetic over the memory model and the attempt log, the tutor has a stored-decomposition fallback, and an AI-off session loads none of the generative code. The system that ships and scores is the system without AI, and AI improves it where it is available and trustworthy.

The data-model decision is a quiet instance of the same discipline. Storing the attempt log as notes rather than a custom table cost a little analytic convenience, but it bought two-way sync for free and kept the most heavily weighted requirement, correct multi-device sync, out of reach of a bug in hand-written sync code. The pre-engineered but unbuilt cache means the convenience can be recovered later without touching sync. Choosing the boring option that protects the load-bearing requirement is, in miniature, the whole posture of the project.

pgrep also sits in a specific relationship to the tool it reuses. A stock spaced-repetition application answers one question well, will you recall this card, and answers it with a strong memory model. pgreg keeps that answer intact and adds the two questions an exam candidate actually asks, can I apply it and what will I score, without disturbing the memory model underneath. The contribution is not a better scheduler, it is the readiness layer and the honesty machinery built on a scheduler that was already good.

Several threads are ready to pick up. The readiness-aware weighting of the selector is designed but not shipped. The Performance model awaits a real attempt history to fit true coefficients, at which

point the same held-out harness that validated the pipeline validates the model itself. The two short AI cutoffs are a generation problem, cutting the malformed-problem rate and tightening fact precision, not a change to the bar. And the calibration measured offline here can be surfaced live as the attempt log grows.

9 Limitations

The central limitation is sample size. `pgrep` is validated by one learner, so several results are mechanism-level rather than population-level. The Performance model is validated on synthetic data, which exercises the pipeline but does not establish the learner’s true coefficients. The interleaving ablation is a simulation under the memory model as ground truth, not a human trial, and its anti-blocking arm is invisible to that model by construction. The simulation also collapses ratings to pass or lapse and assumes the selector’s predicted recall equals the ground-truth recall, which isolates ordering from calibration error but idealizes it. The Readiness check against a real form is a low-sample sanity check. On the AI layer, two absolute quality cutoffs remain just short, card fact precision and the refusal-limited problem batch yield, and closing them is generation work still to do. The literature cited here was assembled from the project’s working notes, and a citation-verification pass for exact figures and identifiers is pending before any external use. None of these undercut the reported numbers, but they bound what those numbers mean.

10 Conclusion

`pgrep` is a Physics GRE preparation system that treats readiness as a measurement problem. It reorders study by exam value inside the engine, it reports three separate and calibrated scores that never confuse recall with transfer, and it generates practice items that cite a named source or refuse, all validated on held-out data and all still working with AI switched off. The evidence supports specific claims and no more: Memory calibrates and beats its baseline on real review logs, the selector beats massed practice everywhere and stock scheduling at realistic budgets, and the AI layer beats every baseline while two absolute quality bars remain in reach. Stated that precisely, the system does what it claims, and it says so honestly where it does not.

Data Availability

The bundled corpus is openly licensed (OpenStax University Physics and the Fitzpatrick texts). Official exam forms used for held-out validation and as a grading ruler are copyright and are never bundled, shipped, indexed, or fed to generation; only numeric score-conversion constants are used, and they stay in a private workspace. The private evaluation workspace (corpus index, gold items, held-out forms, and harness) is not distributed. Aggregate results and the evaluation harness are reproducible from a single command with fixed seeds.

Ethics and AI-Use Disclosure

`pgrep` uses large language models in two places, and this report discloses both. In the product, AI generates study items and tutoring, governed by provenance, gold-set gating, and a leakage firewall,

and disabled by default. In the preparation of this report, an AI assistant helped organize source material and draft prose from the project’s own documentation and recorded results; all numbers trace to the project’s evaluation artifacts, and the author is responsible for the content. The system inherits and preserves the reused engine’s AGPL license and its attribution.

References

- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakcı, Ö., & Mariman, R. (2024). *Generative ai can harm learning* (Tech. Rep.). The Wharton School. (Working paper)
- Bertsch, S., Pesta, B. J., Wiscott, R., & McDaniel, M. A. (2007). The generation effect: A meta-analytic review. *Memory & Cognition*, 35(2), 201–210.
- Brier, G. W. (1950). Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1), 1–3.
- Brown, L. D., Cai, T. T., & DasGupta, A. (2001). Interval estimation for a binomial proportion. *Statistical Science*, 16(2), 101–133.
- Cen, H., Koedinger, K., & Junker, B. (2006). Learning factors analysis: A general method for cognitive model evaluation and improvement. In *Proceedings of the 8th international conference on intelligent tutoring systems* (pp. 164–175).
- Chi, M. T. H., Bassok, M., Lewis, M. W., Reimann, P., & Glaser, R. (1989). Self-explanations: How students study and use examples in learning to solve problems. *Cognitive Science*, 13(2), 145–182.
- Corbett, A. T., & Anderson, J. R. (1994). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4(4), 253–278.
- DeGroot, M. H., & Fienberg, S. E. (1983). The comparison and evaluation of forecasters. *Journal of the Royal Statistical Society: Series D (The Statistician)*, 32(1-2), 12–22.
- Embretson, S. E., & Reise, S. P. (2000). *Item response theory for psychologists*. Lawrence Erlbaum Associates.
- Firth, J., Rivers, I., & Boyle, J. (2021). A systematic review of interleaving as a concept learning strategy. *Review of Education*.
- Galyardt, A., & Goldin, I. (2014). Recent-performance factors analysis. In *Proceedings of the 7th international conference on educational data mining* (pp. 411–412).
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., & Rubin, D. B. (2013). *Bayesian data analysis* (3rd ed.). CRC Press.
- Gjerde, K., et al. (2022). Retrieval practice of principles improves subsequent physics problem solving. *Physical Review Physics Education Research*.
- Gong, Y., Beck, J. E., & Heffernan, N. T. (2011). How to construct more accurate student models: Comparing and optimizing knowledge tracing and performance factor analysis. *International Journal of Artificial Intelligence in Education*, 21(1-2), 27–46.
- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. In *Proceedings of the 34th international conference on machine learning* (pp. 1321–1330).
- Kalyuga, S. (2007). Expertise reversal effect and its implications for learner-tailored instruction. *Educational Psychology Review*, 19(4), 509–539.
- Kapoor, S., & Narayanan, A. (2023). Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9), 100804.
- Kapur, M. (2008). Productive failure. *Cognition and Instruction*, 26(3), 379–424.
- Kapur, M. (2016). Examining productive failure, productive success, unproductive failure, and unproductive success in learning. *Educational Psychologist*, 51(2), 289–299.

- Kornell, N., & Bjork, R. A. (2008). Learning concepts and categories: Is spacing the enemy of induction? *Psychological Science*, 19(6), 585–592.
- Kull, M., Silva Filho, T. M., & Flach, P. (2017). Beta calibration: A well-founded and easily implemented improvement on logistic calibration for binary classifiers. *Proceedings of Machine Learning Research (AISTATS)*, 54, 623–631.
- Le Cam, L. (1960). An approximation theorem for the poisson binomial distribution. *Pacific Journal of Mathematics*, 10(4), 1181–1197.
- Lord, F. M. (1980). *Applications of item response theory to practical testing problems*. Lawrence Erlbaum Associates.
- Metcalfe, J., & Kornell, N. (2005). A region of proximal learning model of study time allocation. *Journal of Memory and Language*, 52(4), 463–477.
- Open Spaced Repetition. (2024a). *anki-revlogs-10k: A dataset of anonymized Anki review logs*. <https://huggingface.co/datasets/open-spaced-repetition/anki-revlogs-10k>.
- Open Spaced Repetition. (2024b). *FSRS: A modern spaced repetition scheduler (Free Spaced Repetition Scheduler)*. <https://github.com/open-spaced-repetition>. (Open-source software and srs-benchmark)
- Pan, S. C., Zung, I., Imundo, M. N., Zhang, X., & Qiu, Y. (2022). User-generated digital flashcards yield better learning than premade flashcards. *Journal of Applied Research in Memory and Cognition*.
- Pavlik, P. I., Cen, H., & Koedinger, K. R. (2009). Performance factors analysis: A new alternative to knowledge tracing. In *Proceedings of the 14th international conference on artificial intelligence in education* (pp. 531–538).
- Peduzzi, P., Concato, J., Kemper, E., Holford, T. R., & Feinstein, A. R. (1996). A simulation study of the number of events per variable in logistic regression analysis. *Journal of Clinical Epidemiology*, 49(12), 1373–1379.
- Piech, C., Bassen, J., Huang, J., Ganguli, S., Sahami, M., Guibas, L., & Sohl-Dickstein, J. (2015). Deep knowledge tracing. In *Advances in neural information processing systems 28* (pp. 505–513).
- Platt, J. C. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in large margin classifiers* (pp. 61–74). MIT Press.
- Renkl, A., & Atkinson, R. K. (2003). Structuring the transition from example study to problem solving in cognitive skill acquisition. *Educational Psychologist*, 38(1), 15–22.
- Roediger, H. L., & Karpicke, J. D. (2006). Test-enhanced learning: Taking memory tests improves long-term retention. *Psychological Science*, 17(3), 249–255.
- Rohrer, D., Dedrick, R. F., & Burgess, K. (2014). The benefit of interleaved mathematics practice is not limited to superficially similar kinds of problems. *Psychonomic Bulletin & Review*, 21(5), 1323–1330.
- Rohrer, D., & Taylor, K. (2007). The shuffling of mathematics problems improves learning. *Instructional Science*, 35, 481–498.
- Samani, J., & Pan, S. C. (2021). Interleaved practice enhances memory and problem-solving ability in undergraduate physics. *npj Science of Learning*, 6(1), 32.
- Sana, F., & Yan, V. X. (2022). Interleaving retrieval practice promotes science learning. *Journal of Applied Research in Memory and Cognition*.
- Settles, B., & Meeder, B. (2016). A trainable spaced repetition model for language learning. In *Proceedings of the 54th annual meeting of the association for computational linguistics* (pp. 1848–1858).
- Slamecka, N. J., & Graf, P. (1978). The generation effect: Delineation of a phenomenon. *Journal of Experimental Psychology: Human Learning and Memory*, 4(6), 592–604.

- Soderstrom, N. C., & Bjork, R. A. (2015). Learning versus performance: An integrative review. *Perspectives on Psychological Science*, 10(2), 176–199.
- Sweller, J. (1988). Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2), 257–285.
- Tabibian, B., Upadhyay, U., De, A., Zarezade, A., Schölkopf, B., & Gomez-Rodriguez, M. (2019). Enhancing human learning via spaced repetition optimization. *Proceedings of the National Academy of Sciences*, 116(10), 3988–3993.
- Taylor, K., & Rohrer, D. (2010). The effects of interleaved practice. *Applied Cognitive Psychology*, 24(6), 837–848.
- Wilson, R. C., Shenhav, A., Straccia, M., & Cohen, J. D. (2019). The eighty five percent rule for optimal learning. *Nature Communications*, 10, 4646.
- Ye, J., Su, J., & Cao, Y. (2022). A stochastic shortest path algorithm for optimizing spaced repetition scheduling. In *Proceedings of the 28th acm sigkdd conference on knowledge discovery and data mining* (pp. 4381–4390).